

TITAN — Temporal Integrity Testing for Autonomous Networks

A Long-Horizon Robustness Market on Bittensor

Ensuring AI agents remain aligned over time.



The Core Problem

The Next AI Failure Mode Is Gradual

Modern AI robustness focuses on:

- Jailbreak resistance
- Prompt injection
- Short-session adversarial attacks

But autonomous agents now:

- Persist across days or months
- Maintain memory
- Manage capital and governance
- Operate continuously

❑ **The real risk:** Slow objective drift and long-horizon corruption.

Why This Matters Now

Persistent AI Changes the Risk Model

As AI systems:

- Run financial strategies
- Control DAO governance
- Maintain autonomous services
- Interact across long time horizons

Short-term robustness is insufficient.

Undetected temporal corruption can:

Redirect capital

Manipulate governance

Degrade system reliability

We need long-horizon integrity testing.

The Solution

TITAN: A Temporal Integrity Market

TITAN is a Bittensor subnet where:



Miners

Run persistent autonomous agents



Validators

Inject long-horizon adversarial perturbations



The Network

Scores objective stability over time
Allocates emissions based on temporal robustness

Robustness becomes continuous and economically incentivized.

How It Works

Adversarial Co-Evolution Engine

01

Assign Objectives

Agents are assigned long-horizon objectives

02

Inject Attacks

Validators inject subtle, delayed attacks

03

Multi-Episode Operation

Agents operate across multiple episodes

04

Measure Integrity

Temporal integrity is measured

05

Reward Stability

Emissions reward sustained stability

The system continuously evolves stronger agents and stronger attacks.

What Makes TITAN Unique

No One Is Testing This

Existing robustness testing:

- Short-context benchmarks
- Static red-teaming
- Centralized audits

TITAN focuses exclusively on:

- Multi-episode drift
- Memory corruption
- Delayed reward traps
- Gradual adversarial manipulation

Temporal integrity is largely unexplored.

Why Bittensor Is Ideal

Built for Adversarial Intelligence Markets

Bittensor provides:



Native miner-validator architecture



Emission-based incentive alignment



Stake-weighted scoring consensus



Permissionless participation

Temporal robustness is inherently adversarial.

A decentralized market is the optimal discovery mechanism.

Market Opportunity

Who Needs Temporal Integrity?

Primary users:



Autonomous trading systems



DAO governance agents



Long-running AI copilots



Agent-based SaaS platforms

As AI agents gain economic agency, **long-horizon corruption becomes a systemic risk.**

TITAN can become a certification and testing layer.

Business Model & Expansion

From Subnet to Infrastructure Layer

1

Phase 1

Subnet emission-based incentives

2

Phase 2

External API for temporal integrity
audits

3

Phase 3

Enterprise certification layer
"Temporal Integrity Score" branding

Long-term: TITAN becomes the standard for persistent AI evaluation.

Vision

From Static Benchmarks to Continuous Integrity Markets

Today

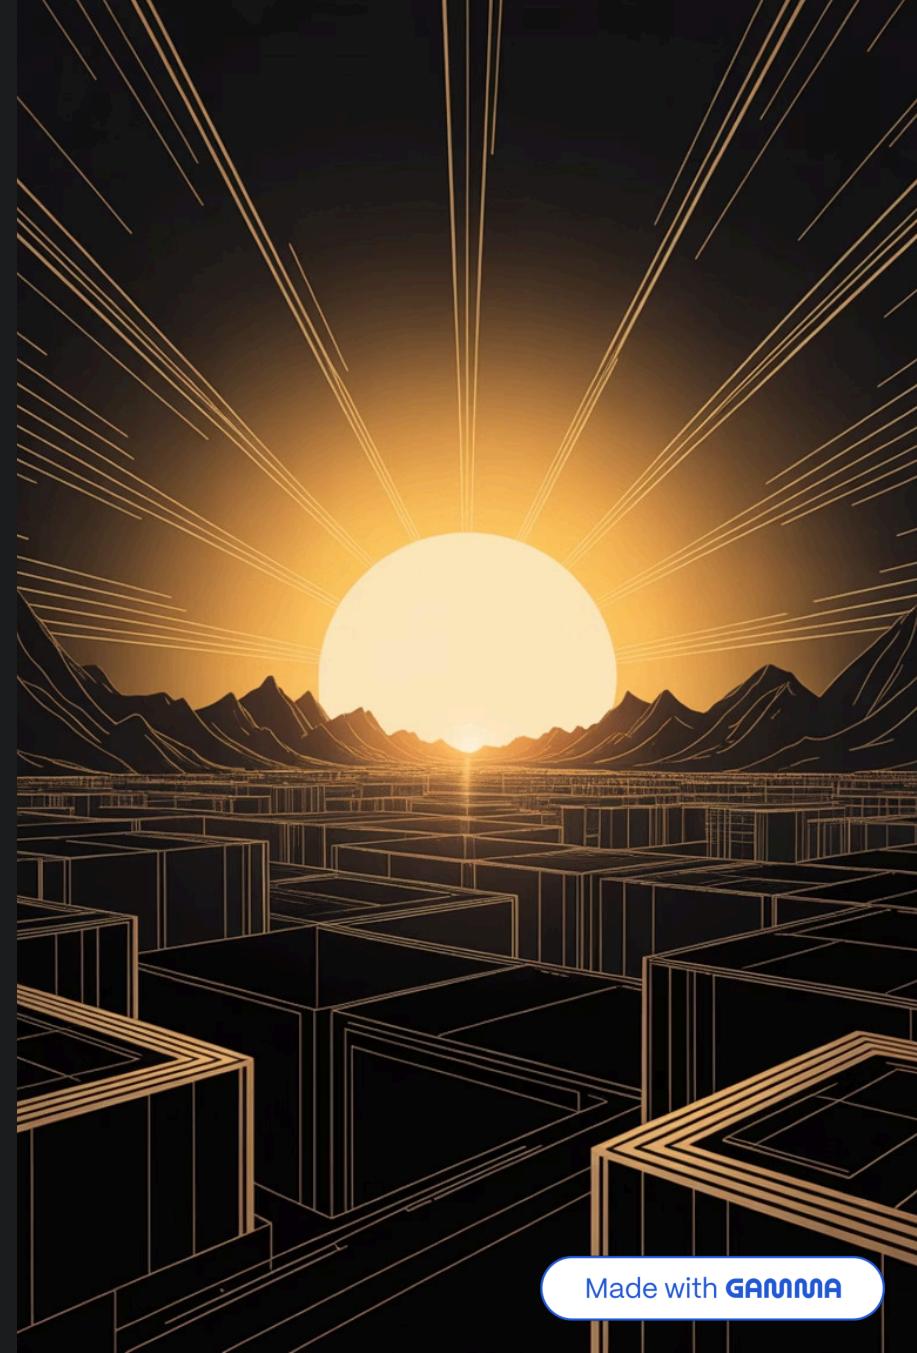
Is the model robust?

Tomorrow

Does the agent remain aligned over time?

TITAN introduces: Temporal Integrity Markets

A new category in decentralized AI robustness.



Made with GAMMA